

COMP7035

Python for Data Analytics and Artificial Intelligence

Pandas

Renjie Wan, Wei Xue

10/30/2022

What we will learn?

<u>Topic</u>	<u>Hours</u>
I. Python Fundamentals A. Program control and logic B. Data types and structures C. Function D. File I/O	12
II. Numerical Computing and Data Visualization Tools and libraries such as A. NumPy B. Matplotlib C. Seaborn	9
III. Exploratory Data Analysis (EDA) with Python Tools and libraries such as A. Pandas B. Sweetviz	9
IV. Artificial Intelligence and Machine Learning with Python Tools and libraries such as A. Keras B. Scikit-learn	9

Contents of Today

- Pandas Series
- Pandas Dataframe: Creation, indexing, slicing


Why do we need Pandas

***Q:** What is the most commonly used **software** for tabular data processing?*

Pandas is a package mainly for working with tabular data.

Main Differences with Numpy Array

- Numpy arrays are designed to contain data of one type (e.g. Int, Float, ...)
- DataFrames can contain different types of data (Int, Float, String, ...)
- Different I/O functions, table operations, time series-specific functionalities...



Numpy Array

The diagram shows a Pandas Dataframe table with the following structure:

	Column Label/ Header	0	1	2	3	4
Index Label		Name	Age	Marks	Grade	Hobby
0	S1	Joe	20	85.10	A	Swimming
1	S2	Nat	21	77.80	B	Reading
2	S3	Harry	19	91.54	A	Music
3	S4	Sam	20	88.78	A	Painting
4	S5	Monica	22	60.55	B	Dancing

Annotations in the diagram:

- Column Index:** Points to the header row (0-4).
- Row Index:** Points to the index column (0-4).
- Column:** Points to a specific column (e.g., Marks).
- Row:** Points to a specific row (e.g., S4).
- Element/ Value/ Entry:** Points to a specific cell (e.g., 88.78).

Pandas Dataframe

Pandas Overview

- Pandas Objects
 - **Series**
 - Dataframe
- Pandas I/O Functions

How to use Pandas

- Just import it!

```
import pandas as pd
```

Series

- Let us start with Series!
- A one-dimensional **labeled** array capable of holding **any mixture** data type
- Axis labels are collectively referred to as the index.
- Think "Series = Vector + labels"
- Create a series: `s = pd.Series(data, index = index)`

```
import numpy as np
import pandas as pd
s = pd.Series(np.random.randn(5), index=['a', 'b', 'c', 'd', 'e'])
print(s)
```

```
a    0.449010
b   -0.391613
c   -1.554223
d    1.054673
e    1.093878
dtype: float64
```


Series

- Creating a series that supports mixed data types

```
import pandas as pd
d = {'a': [0., 0], 'b': {'1':1.}, 'c': 2.}
s = pd.Series(d)
print(s)
```

```
↳ a      [0.0, 0]
   b      {'1': 1.0}
   c              2.0
   dtype: object
```

Series

- Creating a series directly from a dictionary!

```
import pandas as pd
d = {'a': [0., 0], 'b': {'1': 1.}, 'c': 2.}
s = pd.Series(d)
print(s)
```

```
a    [0.0, 0]
b    {'1': 1.0}
c          2.0
dtype: object
```

Index is constructed as sorted keys

- If you want to specify the order of index:

```
s = pd.Series(d, index=['c', 'b', 'a'])
print(s)
```

```
c          2.0
b    {'1': 1.0}
a    [0.0, 0]
dtype: object
```

Series

- Convert the created Series to other data types

```
import pandas as pd
import numpy as np

dict = {'a': 100, 'b': 200, 'c':300, 'd':400, 'e':500}
print("Original dictionary:")
print(dict)
s = pd.Series(dict)
print(s)
s_list= s.to_list()
s_dict = s.to_dict()
print(s_list)
print(s_dict)
```

```
Original dictionary:
{'a': 100, 'b': 200, 'c': 300, 'd': 400, 'e': 500}
a    100
b    200
c    300
d    400
e    500
dtype: int64
[100, 200, 300, 400, 500]
{'a': 100, 'b': 200, 'c': 300, 'd': 400, 'e': 500}
```

Use `x.to_list()` to convert to list
Use `x.to_dict()` to convert to the dict

Series

- Exercise

1. Create a 5-D random numpy list **var_list**
2. use “uuid” to generate 5 random keys (use `str(uuid.uuid4())`), and store them into a list **key_list**
3. Create a dictionary **dict** from **var_list** and **key_list**
4. Create a Pandas Series from
 - a) **var_list**
 - b) **var_list** and **key_list**
 - c) **dict**
5. Convert the Series back to the list and dictionary

Series

- Indexing: Just like you would for NumPy arrays/python lists!

```
import numpy as np
import pandas as pd
s = pd.Series(np.random.rand(5), index=['b', 'a', 'c', 'd', 'e'])
print(s)
print('s[0]={}'.format(s[0]))
print('s["a"]={}'.format(s['a']))
```

```
b    0.324696
a    0.121395
c    0.229814
d    0.555360
e    0.764731
dtype: float64
s[0]=0.3246963251015317
s["a"]=0.12139490883771598
```

Series

- Slicing: Also similar to python lists

```
import numpy as np
import pandas as pd
end_string = '\n' + '-'*50 + '\n'
s = pd.Series(np.random.rand(5), index=['b', 'a', 'c', 'd', 'e'])
print(s, end=end_string)
print(s[:2], end=end_string)
print(s['a':'d'], end=end_string)
```

```
b    0.661580
a    0.235360
c    0.616469
d    0.326972
e    0.132890
dtype: float64
```

```
b    0.66158
a    0.23536
dtype: float64
```

```
a    0.235360
c    0.616469
d    0.326972
dtype: float64
```

Note what elements are selected

Series

- Slicing: picking elements under certain conditions

```
import numpy as np
import pandas as pd
end_string = '\n' + '-'*50 + '\n'
s = pd.Series(np.random.rand(5), index=['b', 'a', 'c', 'd', 'e'])
print(s, end=end_string)
print(s[s>0.5], end=end_string)
```

```
b    0.682803
a    0.020967
c    0.846114
d    0.676687
e    0.279662
dtype: float64
```

```
b    0.682803
c    0.846114
d    0.676687
dtype: float64
```

Series

- Assign new values and indexes

```
import numpy as np
import pandas as pd
end_string = '\n' + '-'*50 + '\n'
s = pd.Series(np.random.rand(5), index=['b', 'a', 'c', 'd', 'e'])
print(s,end=end_string)
s['a']=0
s['f'] = 'test'
print(s,end=end_string)
```

b 0.564049
a 0.387429
c 0.602721
d 0.327507
e 0.386777
dtype: float64

b 0.564049
a 0.0
c 0.602721
d 0.327507
e 0.386777
f test
dtype: object

Series

- Operations
 - Get the element

```
import numpy as np
import pandas as pd
end_string = '\n' + '-'*50 + '\n'
s = pd.Series(np.random.rand(5), index=['b', 'a', 'c', 'd', 'e'])
print(s, end = end_string)
print('f' in s, end = end_string) # check for index label
print(s.get('f', None), end = end_string) # get item with index 'f' - if no such item return None
print(s.get('e', None), end = end_string)
```

```
b    0.628035
a    0.261522
c    0.354718
d    0.345538
e    0.987133
dtype: float64
```

```
False
```

```
None
```

```
0.9871332517316361
```

Note what value is returned

Series

- Exercise

1. Create a 5-D random numpy list **var_list**
 2. use “uuid” to generate 5 random keys (use `str(uuid.uuid4())`), and store them into a list **key_list**
 3. Create a dictionary **dict** from **var_list** and **key_list**
 4. Create a Pandas Series from
 - a) **var_list**
 - b) **var_list** and **key_list**
 - c) **dict**
 5. Convert the Series back to the list and dictionary
-
6. Find out the elements larger than zero
 7. Calculate the proportion of positive elements in the Series

Series

- Operations
 - Math calculations. Numpy operations can be applied to the Series.

```
import numpy as np
import pandas as pd
end_string = '\n' + '-'*50 + '\n'
s = pd.Series(np.random.rand(5), index=['b', 'a', 'c', 'd', 'e'])
print(s, end=end_string)
print(np.exp(s), end=end_string)
```

```
↳ b    0.685840
   a    0.315714
   c    0.994083
   d    0.413190
   e    0.220100
dtype: float64
```

```
-----
b    1.985440
a    1.371238
c    2.702245
d    1.511633
e    1.246201
dtype: float64
-----
```

Series

- Attributes
 - Get the index, value and shape

```
import numpy as np
import pandas as pd
end_string = '\n' + '-'*50 + '\n'
s = pd.Series(np.random.rand(5), index=['b', 'a', 'c', 'd', 'e'])
print(s, end=end_string)
print(s.index, end=end_string)
print(s.values, end=end_string)
print(s.shape, end=end_string)
```

```
b    0.852044
a    0.885222
c    0.509656
d    0.414917
e    0.927696
dtype: float64

-----
Index(['b', 'a', 'c', 'd', 'e'], dtype='object')
-----
[0.85204352 0.88522209 0.50965594 0.41491689 0.92769556]
-----
(5,)
```

Series

- Iteration

```
import numpy as np
import pandas as pd
end_string = '\n' + '-'*50 + '\n'
s = pd.Series(np.random.rand(5), index=['b', 'a', 'c', 'd', 'e'])
print(s, end=end_string)
for idx, val in s.iteritems():
    print(idx, val)
```

```
b 0.354212
a 0.680097
c 0.878909
d 0.322548
e 0.997003
dtype: float64
```

```
-----
b 0.35421193446975086
a 0.6800973399869802
c 0.8789085260364188
d 0.3225477405798104
e 0.9970027318704187
```

Series

- Exercise

1. Create a 5-D random numpy list **var_list**
 2. use “uuid” to generate 5 random keys (use `str(uuid.uuid4())`), and store them into a list **key_list**
 3. Create a dictionary **dict** from **var_list** and **key_list**
 4. Create a Pandas Series **ps** from
 - a) **var_list**
 - b) **var_list** and **key_list**
 - c) **dict**
 5. Convert the Series back to the list and dictionary
-
6. Find out the elements larger than zero
 7. Calculate the proportion of positive elements in the Series
-
8. Write down as many ways of forming a list that contains the values of Series elements
 9. Calculate the proportion of elements that are larger than the mean value of the Series

Pandas Overview

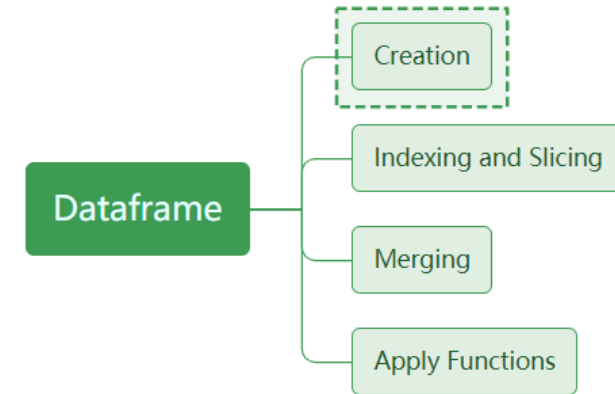
- Pandas Objects
 - Series
 - **Dataframe**
- Pandas I/O Functions

Dataframe

- A two-dimensional **labeled** data structure capable of holding **any mixture** data type
- **Think Dataframe as spreadsheets**
- Create a series: `df = pd.DataFrame(data, index = index, columns = columns)`

Dataframe

- Creating a Dataframe
 - From dict of series or dicts
 - **From dict of series**
 - **From dicts**




```
# Create a dataframe from a dictionary of series
d = {'two': pd.Series([1,2], index = ['a','e']),
     'one': pd.Series(list(range(4)), index = ['b','a', 'c', 'd'])}
df = pd.DataFrame(d)
print(df)
```

```
two one
a  1.0  1.0
b  NaN  0.0
c  NaN  2.0
d  NaN  3.0
e  2.0  NaN
```

Dataframe

- Creating a Dataframe
 - From dict of series or dicts
 - **From dict of series**
 - **From dicts**

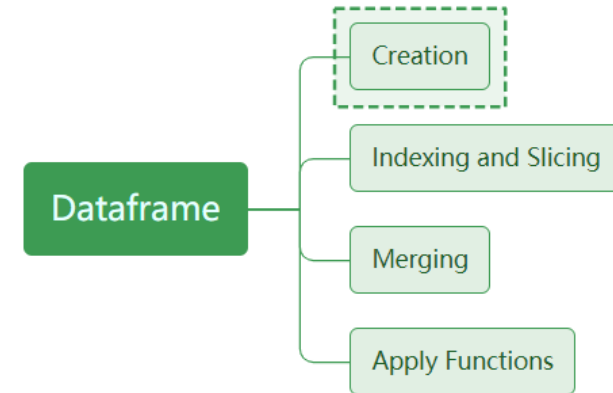


If there are any nested dicts, these will be first converted to Series.

```
d= {'one': {'a': 1, 'b': 2, 'c': 3},  
     'two': pd.Series(list(range(4)), index = ['a', 'b', 'c', 'd'])}  
df = pd.DataFrame(d)  
print(df)
```

```
one two  
a  1.0  0  
b  2.0  1  
c  3.0  2  
d  NaN  3
```

Dataframe



- Creating a Dataframe

- **From dict of ndarray / lists**

- The ndarrays must all be the same length.
- If an index is passed, it must clearly also be the same length as the arrays.
- If no index is passed, the result will be range(len_array)

```
[10] d = {'one' : [1., 2., 3., 4], 'two' : [4., 3., 2., 1.]}
      pd.DataFrame(d)
```

	one	two
0	1.0	4.0
1	2.0	3.0
2	3.0	2.0
3	4.0	1.0

```
d = {'one' : [1., 2., 3., 4], 'two' : [4., 3., 2., 1.]}
pd.DataFrame(d, index=['a', 'b', 'd', 'e'])
```

	one	two
a	1.0	4.0
b	2.0	3.0
d	3.0	2.0
e	4.0	1.0

Dataframe

- Creating a Dataframe
 - From a list of dicts

```
[13] data = []
      for i in range(5):
          data += [ {'Column' + str(j):np.random.randint(100) for j in range(5)} ]
          # dictionary comprehension!
```

```
data[:5]
```

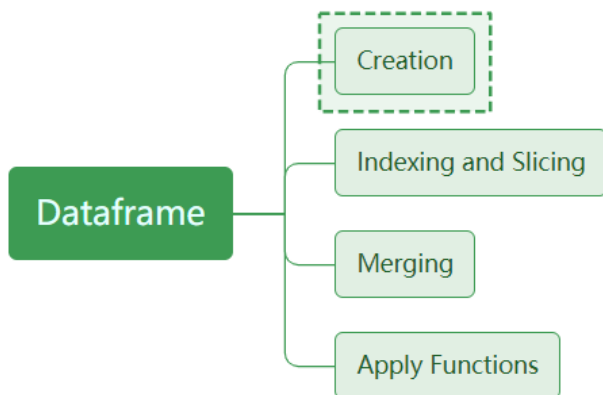
```
[{'Column0': 8, 'Column1': 35, 'Column2': 71, 'Column3': 52, 'Column4': 37},
 {'Column0': 5, 'Column1': 74, 'Column2': 17, 'Column3': 75, 'Column4': 77},
 {'Column0': 83, 'Column1': 76, 'Column2': 7, 'Column3': 79, 'Column4': 55},
 {'Column0': 42, 'Column1': 69, 'Column2': 84, 'Column3': 71, 'Column4': 49},
 {'Column0': 29, 'Column1': 21, 'Column2': 89, 'Column3': 76, 'Column4': 82}]
```



```
df = pd.DataFrame(data)
print(df, end = end_string)
df = pd.DataFrame(data, columns = ['Column0', 'Column1'])
print(df, end = end_string)
```

	Column0	Column1	Column2	Column3	Column4
0	8	35	71	52	37
1	5	74	17	75	77
2	83	76	7	79	55
3	42	69	84	71	49
4	29	21	89	76	82

	Column0	Column1
0	8	35
1	5	74
2	83	76
3	42	69
4	29	21



Series

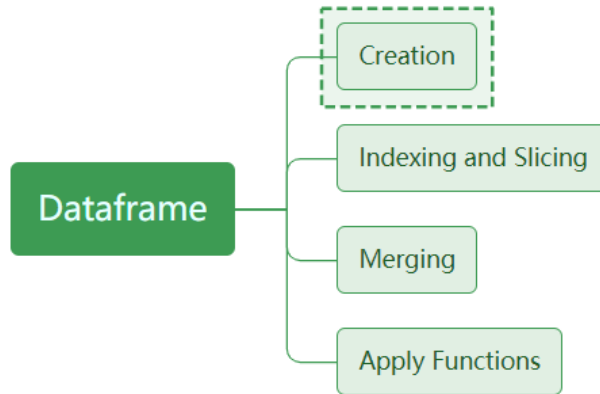
- Exercise

1. Write codes to create a random dict **x** which has 5 random keys and each key corresponds to a 6-D numpy array
2. Create a pandas dataframe using **x**
3. Create a pandas dataframe using a subset of **x**, in the subset of **x**, only keys that start with a digit are chosen.

Dataframe

- Attributes

- `df.index` : the row index of `df`
- `df.columns` : the columns of `df`
- `df.shape` : the shape of the `df`
- `df.values` : numpy array of values



```
df = pd.DataFrame(data, columns = ['Column0', 'Column1', 'Column2', 'Column3'], index=['a', 'b', 'c', 'd', 'e'])
print(df, end = end_string)
print(df.index, end = end_string)
print(df.columns, end = end_string)
print(df.shape, end = end_string)
print(df.values, end = end_string)
```

	Column0	Column1	Column2	Column3
a	8	35	71	52
b	5	74	17	75
c	83	76	7	79
d	42	69	84	71
e	29	21	89	76

```
Index(['a', 'b', 'c', 'd', 'e'], dtype='object')
```

```
Index(['Column0', 'Column1', 'Column2', 'Column3'], dtype='object')
```

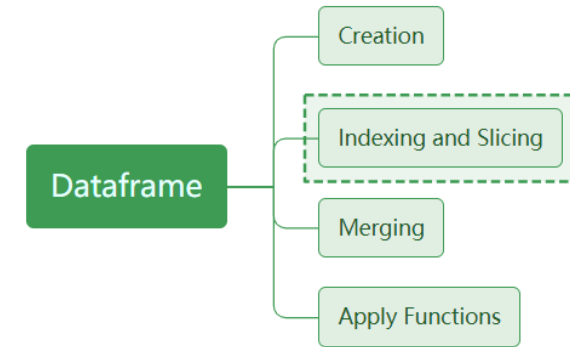
```
(5, 4)
```

```
[ [ 8 35 71 52]
  [ 5 74 17 75]
  [83 76  7 79]
  [42 69 84 71]
  [29 21 89 76]]
```

Dataframe


- Indexing and Slicing

3 methods [], iloc, loc



Operation	Syntax	Result
Select Column	df[col]	Series
Select Row by Label	df.loc[label]	Series
Select Row by Integer Location	df.iloc[idx]	Series
Select Columns	df[col_list]	DataFrame
Slice rows	df[5:10]	DataFrame
Select rows by boolean	df[mask]	DataFrame

Dataframe



- Simplest form of Indexing: []

Operation	Syntax	Result
Select Column	df[col]	Series
Select Columns	df[col_list]	DataFrame
Slice rows	df[5:10]	DataFrame
Select rows by boolean	df[mask]	DataFrame

```

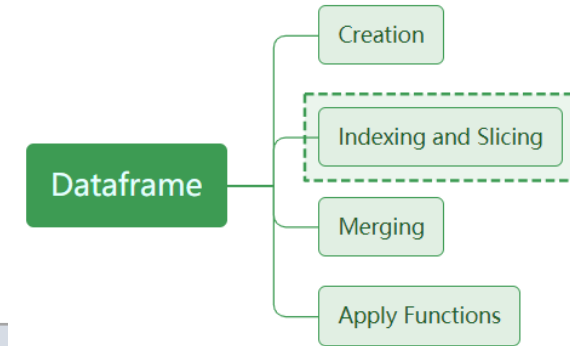
# Lets create a data frame
pd.options.display.max_rows = 4
dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
df
  
```

	A	B	C	D
2000-01-01	-0.162525	-0.784332	0.556824	-0.564598
2000-01-02	-0.203106	0.978618	-1.308789	1.738307
...
2000-01-07	0.020128	0.132962	0.529890	1.230876
2000-01-08	-1.264452	0.017298	1.405524	-0.097113

8 rows x 4 columns

Let us create a dataframe first

Dataframe



- Simplest form of Indexing: []

Operation	Syntax	Result
Select Column	df[col]	Series
Select Columns	df[col_list]	DataFrame
Slice rows	df[5:10]	DataFrame
Select rows by boolean	df[mask]	DataFrame

```

▶ # Lets create a data frame
pd.options.display.max_rows = 4
dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
# column 'A'
df['A']

```

```

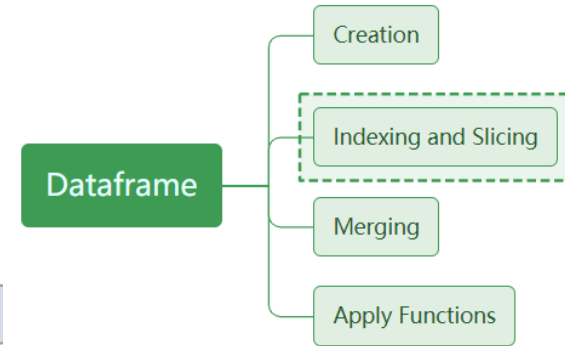
2000-01-01    1.258281
2000-01-02   -1.753082
...
2000-01-07   -1.363351
2000-01-08    1.506044
Freq: D, Name: A, Length: 8, dtype: float64

```

Dataframe

- Simplest form of Indexing: []

Operation	Syntax	Result
Select Column	df[col]	Series
Select Columns	df[col_list]	DataFrame
Slice rows	df[5:10]	DataFrame
Select rows by boolean	df[mask]	DataFrame



```

# Lets create a data frame
pd.options.display.max_rows = 4
dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
# column 'A' and 'C'
df[['A', 'C']]
  
```

	A	C
2000-01-01	0.426504	-0.994245
2000-01-02	0.124058	1.185978
...
2000-01-07	-0.142703	-0.387663
2000-01-08	0.307563	-0.460706

8 rows x 2 columns

Dataframe

- Simplest form of Indexing: []

Operation	Syntax	Result
Select Column	df[col]	Series
Select Columns	df[col_list]	DataFrame
Slice rows	df[5:10]	DataFrame
Select rows by boolean	df[mask]	DataFrame

Dataframe

Creation

Indexing and Slicing

Merging

Apply Functions

```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)
print(df[2:5], end=end_string)

```

```

      A      B      C      D
2000-01-01 -1.277544 -0.442661 -0.216973 -0.884319
2000-01-02  0.019020  0.032021 -2.701674  0.635646
2000-01-03 -0.992079  0.829751  0.373975 -1.756382
2000-01-04  1.164141 -0.718194  0.421115 -1.192541
2000-01-05  0.633056  0.008495  0.458766  0.216892
2000-01-06 -0.148220  0.761876 -0.591145  0.318806
2000-01-07  0.749281 -2.055792 -0.419089 -0.519670
2000-01-08  0.209152 -0.893017  0.645123  0.108354

```

```

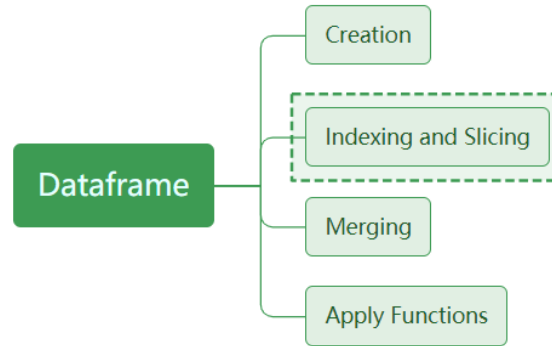
      A      B      C      D
2000-01-03 -0.992079  0.829751  0.373975 -1.756382
2000-01-04  1.164141 -0.718194  0.421115 -1.192541
2000-01-05  0.633056  0.008495  0.458766  0.216892

```

Dataframe

- Simplest form of Indexing: []

Operation	Syntax	Result
Select Column	df[col]	Series
Select Columns	df[col_list]	DataFrame
Slice rows	df[5:10]	DataFrame
Select rows by boolean	df[mask]	DataFrame



```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)
print(df[df['A'] > df['B']], end=end_string)
  
```

Here, a boolean mask is defined by some conditions.

```

➞
      A      B      C      D
2000-01-01  0.379490  1.353836 -0.634997 -0.066652
2000-01-02  0.635418 -1.348018 -0.310449  0.278415
2000-01-03 -0.582322 -0.539813  0.785476  1.236848
2000-01-04 -0.130398  1.050291 -0.832720 -1.692569
2000-01-05  1.040468 -1.033498  0.948710 -0.085250
2000-01-06  1.307932  2.202969  1.640022 -2.448868
2000-01-07 -0.203380 -1.231041  0.785852 -0.141751
2000-01-08 -1.234723 -0.745927  1.268844 -0.412678
-----
      A      B      C      D
2000-01-02  0.635418 -1.348018 -0.310449  0.278415
2000-01-05  1.040468 -1.033498  0.948710 -0.085250
2000-01-07 -0.203380 -1.231041  0.785852 -0.141751
-----
  
```

Note the values of A and B columns

Series


- Exercise

1. Write codes to create a random dict **x** which has 5 random keys and each key corresponds to a 6-D numpy array
 2. Create a pandas dataframe using **x**
 3. Create a pandas dataframe using a subset of **x**, in the subset of **x**, only keys that start with a digit are chosen.
-
4. Create a new pandas dataframe using the codes in the previous slide.
 5. Select rows whose attribute A is smaller than the mean of attribute C
 6. Can you select the column B and C using [] indexing? Try it out and see what happens.

Dataframe

- Selecting by label .loc (string based)

Operation	Syntax	Result
Select Row by Label	df.loc[label]	Series or dataframe



Allowed inputs:

1. A single label
2. A list of labels
3. A boolean array

```

▶ dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)
print(df.loc['2000-01-01'], end=end_string)

```

```

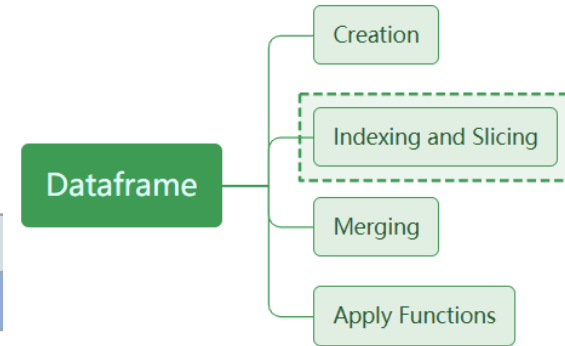
↵
      A      B      C      D
2000-01-01 -1.476048  0.945996  0.017978  0.468844
2000-01-02 -0.048315 -0.231605  1.783954 -0.051215
2000-01-03 -2.247547 -0.701498 -0.240208 -1.853042
2000-01-04  0.752344  0.127446  0.426287  1.159992
2000-01-05  1.379715  0.474799  0.448713 -0.573342
2000-01-06 -0.141288 -0.327275 -0.741280  0.101982
2000-01-07 -0.529335 -0.884688  0.952924 -0.638781
2000-01-08  0.192356 -0.347582 -1.761751  0.221353
-----
A    -1.476048
B      0.945996
C      0.017978
D      0.468844
Name: 2000-01-01 00:00:00, dtype: float64
-----

```

Dataframe

- Selecting by label .loc (string based)

Operation	Syntax	Result
Select Row by Label	df.loc[label]	Series or dataframe



Allowed inputs:

1. A single label
2. A list of labels
3. A boolean array

```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)
print(df.loc[:, 'A'], end=end_string)
  
```


```

┌───┐
      A      B      C      D
2000-01-01  1.335074  0.245156  1.503536 -0.399030
2000-01-02 -0.607126  0.586251  0.123652  0.974062
2000-01-03 -0.823209 -0.030745 -0.712795 -0.361908
2000-01-04  0.903869  0.455098  0.615823  2.187766
2000-01-05 -0.508980 -0.563408 -1.990079 -0.589906
2000-01-06 -0.227697  0.037874  0.528425 -0.285479
2000-01-07 -0.420371  0.890651 -0.578147  1.272716
2000-01-08 -1.927446  0.506257 -0.108151  1.600303
-----
2000-01-01    1.335074
2000-01-02   -0.607126
2000-01-03   -0.823209
2000-01-04    0.903869
2000-01-05   -0.508980
2000-01-06   -0.227697
2000-01-07   -0.420371
2000-01-08   -1.927446
Freq: D, Name: A, dtype: float64
  
```

Dataframe

- Selecting by label .loc (string based)

Operation	Syntax	Result
Select Row by Label	df.loc[label]	Series or dataframe



Allowed inputs:

1. A single label
2. A list of labels
3. A boolean array

```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)
print(df.loc['2000-01-01':'2000-01-03', 'A':'C'], end=end_string)
  
```

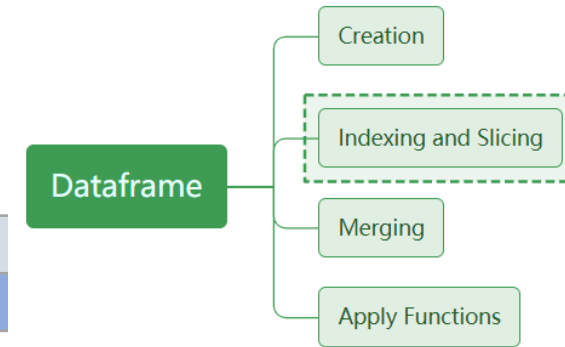
```

  A      B      C      D
2000-01-01  2.242905  0.603600 -0.605938  1.884147
2000-01-02 -0.374925  0.200928  0.793096  1.440772
2000-01-03 -0.878849  1.156089 -0.263019 -2.007993
2000-01-04 -1.672786 -0.675574  2.878655  1.160563
2000-01-05  0.161643 -2.200404 -2.285876 -0.075032
2000-01-06 -2.518145  1.210010 -1.792641  0.167813
2000-01-07  1.256607  0.615687 -0.402351 -0.589810
2000-01-08  0.343399 -0.623055  1.679456 -0.675518
-----
               A      B      C
2000-01-01  2.242905  0.603600 -0.605938
2000-01-02 -0.374925  0.200928  0.793096
2000-01-03 -0.878849  1.156089 -0.263019
  
```


Dataframe

- Selecting by label .loc (string based)

Operation	Syntax	Result
Select Row by Label	df.loc[label]	Series or dataFrame



Allowed inputs:

1. A single label
2. A list of labels
3. A boolean array

```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)
print(df.loc[:, df.loc['2000-01-01'] > 0], end=end_string)
  
```

```

  A      B      C      D
2000-01-01 -0.566362  1.381827  0.303849 -0.699103
2000-01-02  1.056164 -0.921311 -0.431884  1.174159
2000-01-03  0.584125  0.140604  1.197306 -0.713242
2000-01-04 -0.441814 -0.580099  0.538301 -2.244958
2000-01-05 -1.942451  0.195014  0.269126 -1.058965
2000-01-06 -0.972741  0.968222  1.115717 -0.147523
2000-01-07  0.503898  0.380396 -0.903563 -1.037288
2000-01-08 -0.732533 -0.013822  1.121445 -0.002982
  
```

```

      B      C
2000-01-01  1.381827  0.303849
2000-01-02 -0.921311 -0.431884
2000-01-03  0.140604  1.197306
2000-01-04 -0.580099  0.538301
2000-01-05  0.195014  0.269126
2000-01-06  0.968222  1.115717
2000-01-07  0.380396 -0.903563
2000-01-08 -0.013822  1.121445
  
```

Note the values of selected columns
for the row 2000-01-01

Series

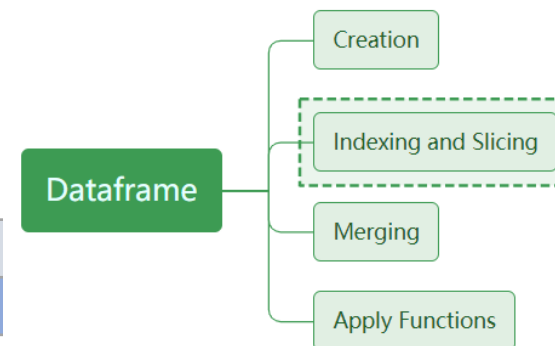
- Exercise

1. Write codes to create a random dict **x** which has 5 random keys and each key corresponds to a 6-D numpy array
 2. Create a pandas dataframe using **x**
 3. Create a pandas dataframe using a subset of **x**, in the subset of **x**, only keys that start with a digit are chosen.
-
4. Create a new pandas dataframe using the codes in the previous slide.
 5. Select rows whose attribute A is smaller than the mean of attribute C
 6. Can you select the column B and C using [] indexing? Try it out and see what happens.
-
7. Now try to select the column B and C using the newly learned method.

Dataframe

- Selecting by position .iloc (index based)

Operation	Syntax	Result
Select Row by Integer Location	df.iloc[idx]	Series/Dataframe



Allowed inputs:

1. An integer
2. A list of integers
3. A slice
4. A boolean array

```

▶ dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)
print(df.iloc[3], end=end_string)

```

```


A      B      C      D
2000-01-01  0.190067  0.495738  0.408074 -0.842692
2000-01-02 -0.709675 -1.498102  0.416754 -1.989156
2000-01-03  0.911487  0.029719 -0.799064 -0.356679
2000-01-04  1.172371 -0.222231 -0.277088  1.550702
2000-01-05 -0.537345  1.798668 -0.629828 -1.349256
2000-01-06 -1.253146  0.057198  0.612258 -0.199975
2000-01-07 -0.255768 -0.788770  1.092821  0.662147
2000-01-08 -0.172003 -1.277385 -1.931545  0.201848
-----
A      1.172371
B     -0.222231
C     -0.277088
D      1.550702
Name: 2000-01-04 00:00:00, dtype: float64
-----

```

Dataframe

- Selecting by position .iloc (index based)

Operation	Syntax	Result
Select Row by Integer Location	df.iloc[idx]	Series/Dataframe



Allowed inputs:

1. An integer
2. A list of integers
3. A slice
4. A boolean array

```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)
print(df.iloc[[0, 2], [1, 2]], end=end_string)
  
```

```

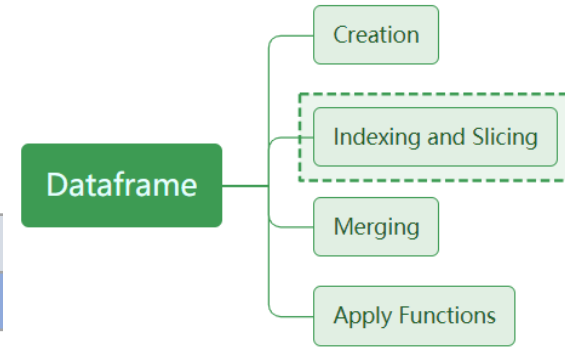
      A         B         C         D
2000-01-01  0.699167  0.179175  0.369038  0.487041
2000-01-02 -0.925355 -0.774855 -0.334079 -1.125520
2000-01-03 -1.816214 -0.583109  0.597376  0.652739
2000-01-04  0.198655  1.590933 -0.280309 -0.392774
2000-01-05 -0.491652  0.476999 -2.199367 -0.716448
2000-01-06 -1.061156 -0.662827  0.974857 -0.230955
2000-01-07  0.000544 -0.720861 -1.032508  0.228881
2000-01-08  0.077399  2.617564  0.799700 -0.555486

-----
              B         C
2000-01-01  0.179175  0.369038
2000-01-03 -0.583109  0.597376
  
```

Dataframe

- Selecting by position .iloc (index based)

Operation	Syntax	Result
Select Row by Integer Location	df.iloc[idx]	Series/Dataframe



Allowed inputs:

1. An integer
2. A list of integers
3. A slice
4. A boolean array

```

▶ dates = pd.date_range('1/1/2000', periods=8)
  df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
  end_string = '\n' + '-'*50 + '\n'
  print(df, end=end_string)
  print(df.iloc[:2, 2:], end=end_string)
  
```

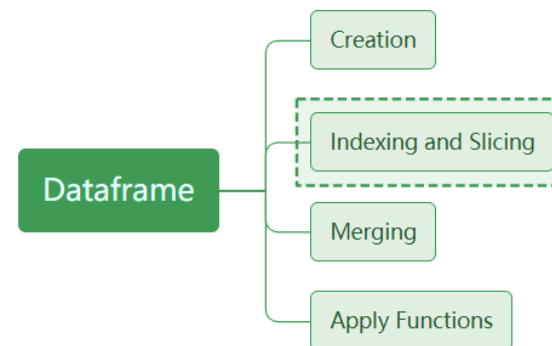
```

              A          B          C          D
2000-01-01 -0.543968 -0.872015 -0.073505  0.075670
2000-01-02  0.025429 -0.480559 -1.193146  0.154151
2000-01-03  0.565404  0.461674 -2.500733 -0.460552
2000-01-04 -0.560048  1.579918  0.119610  0.949158
2000-01-05 -1.230365 -1.058706  0.728648  0.942845
2000-01-06 -0.746050 -0.938077 -0.636031 -1.273335
2000-01-07 -0.508007  0.806813  2.299999  0.899051
2000-01-08  0.543415  0.172509  2.424232  0.505163
-----
              C          D
2000-01-01 -0.073505  0.075670
2000-01-02 -1.193146  0.154151
-----
  
```

Dataframe

- Selecting by position `.iloc` (index based)

Operation	Syntax	Result
Select Row by Integer Location	<code>df.iloc[idx]</code>	Series/Dataframe



Allowed inputs:

1. An integer
2. A list of integers
3. A slice
4. A boolean array

```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)

boolean_mask = df.iloc[:, 1] > 0.0
print(boolean_mask.values, end=end_string)
print(df.iloc[boolean_mask.values, :], end=end_string)
  
```

Selecting the dates with positive values in column B

```

A      B      C      D
2000-01-01  1.289575 -0.747892  0.199428  0.309966
2000-01-02  0.117910  0.898768 -0.248015 -0.334323
2000-01-03  0.596167 -0.009834 -1.062685 -0.225856
2000-01-04 -0.438745  0.274453  2.070763  1.005517
2000-01-05 -2.085728  0.108528  0.295392  0.334122
2000-01-06 -0.660009  1.318652  0.277901 -0.323657
2000-01-07 -2.143342 -0.406465 -1.494151 -0.016955
2000-01-08 -2.267976 -0.964705 -1.692748  1.739006


[False  True False  True  True  True False False]

A      B      C      D
2000-01-02  0.117910  0.898768 -0.248015 -0.334323
2000-01-04 -0.438745  0.274453  2.070763  1.005517
2000-01-05 -2.085728  0.108528  0.295392  0.334122
2000-01-06 -0.660009  1.318652  0.277901 -0.323657
  
```

Dataframe

- Selecting by position `.iloc` (index based)

Operation	Syntax	Result
Select Row by Integer Location	<code>df.iloc[idx]</code>	Series/Dataframe



Allowed inputs:

1. An integer
2. A list of integers
3. A slice
4. A boolean array

```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)

boolean_mask = df.iloc[:, 1] > 0.0
print(boolean_mask.values, end=end_string)
print(df.iloc[boolean_mask.values, :], end=end_string)
  
```

```

      A         B         C         D
2000-01-01  1.289575 -0.747892  0.199428  0.309966
2000-01-02  0.117910  0.898768 -0.248015 -0.334323
2000-01-03  0.596167 -0.009834 -1.062685 -0.225856
2000-01-04 -0.438745  0.274453  2.070763  1.005517
2000-01-05 -2.085728  0.108528  0.295392  0.334122
2000-01-06 -0.660009  1.318652  0.277901 -0.323657
2000-01-07 -2.143342 -0.406465 -1.494151 -0.016955
2000-01-08 -2.267976 -0.964705 -1.692748  1.739006

[False  True False  True  True  True False False]

      A         B         C         D
2000-01-02  0.117910  0.898768 -0.248015 -0.334323
2000-01-04 -0.438745  0.274453  2.070763  1.005517
2000-01-05 -2.085728  0.108528  0.295392  0.334122
2000-01-06 -0.660009  1.318652  0.277901 -0.323657
  
```

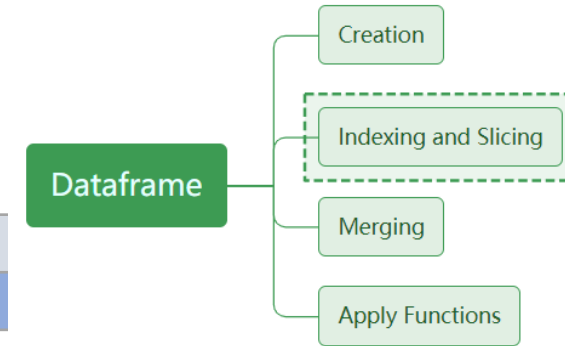
Selecting the dates with positive values in column B

Recap: How to achieve this by using `.loc`?

Dataframe

- Selecting by position .iloc (index based)

Operation	Syntax	Result
Select Row by Integer Location	df.iloc[idx]	Series/Dataframe



Allowed inputs:

1. An integer
2. A list of integers
3. A slice
4. A boolean array

```

dates = pd.date_range('1/1/2000', periods=8)
df = pd.DataFrame(np.random.randn(8, 4), index=dates, columns=['A', 'B', 'C', 'D'])
end_string = '\n' + '-'*50 + '\n'
print(df, end=end_string)

print(df.loc[df.loc[:, 'B']>0, :], end=end_string)
  
```

```

      A         B         C         D
2000-01-01  1.478031 -0.182063 -0.863153 -0.782226
2000-01-02  1.390328  0.830639 -0.721648  1.468268
2000-01-03 -0.240887 -0.627681 -0.478662 -0.585520
2000-01-04  1.204037 -0.293523 -0.056818  0.413413
2000-01-05 -1.302853 -0.644060  0.902543 -0.070009
2000-01-06  0.539073  0.790697  0.734030 -0.324839
2000-01-07 -0.869246  0.052510 -0.768528  1.598001
2000-01-08 -1.080327 -0.238485 -1.414379  0.021658
-----
      A         B         C         D
2000-01-02  1.390328  0.830639 -0.721648  1.468268
2000-01-06  0.539073  0.790697  0.734030 -0.324839
2000-01-07 -0.869246  0.052510 -0.768528  1.598001
-----
  
```

Selecting the dates with positive values in column B

Recap: How to achieve this by using .loc?

Series

- Exercise

1. Write codes to create a random dict **x** which has 5 random keys and each key corresponds to a 6-D numpy array
 2. Create a pandas dataframe using **x**
 3. Create a pandas dataframe using a subset of **x**, in the subset of **x**, only keys that start with a digit are chosen.
-
4. Create a new pandas dataframe using the codes in the previous slide.
 5. Select rows whose attribute A is smaller than the mean of attribute C
 6. Can you select the column B and C using [] indexing? Try it out and see what happens.
-
7. Now try to select the column B and C using the newly learned method.
-
8. Select the upper right half of the dataframe to create a new dataframe “**ur_df**”, find out the largest value in the last row of the **ur_df**.